

# Traffic Predictions

Data 621 Final

Tommy Jenkins, Violeta Stoyanova, Todd Weigel, Peter Kowalchuk, Eleanor R-Secoquian,  
Anthony Pagan

2019/12/04

- Utility functions
- OVERVIEW
- Objective:
- DATA EXPLORATION
  - Load data
    - View Data
  - Data Summary
    - Missing values
  - Reformating features
- Visualizations
  - Trip duration vs pickup and dropoff datetime and location using 20% of sampled data
  - More visualizations
- An excursion into classification
- Model, correlation
- Modeling
  - Sample Model 1 Poisson (No store\_and\_fwd\_flag, id or vendor\_id)
  - Sample Model 2 Gaussian (No id or vendor\_id)
  - Sample Model 3 Negative Binomial (No id)
- CONCLUSION
- APPENDIX
- References:

## Utility functions

From [R Cookbooks](#) to create multi-panel plots.

## OVERVIEW

In this competition, Kaggle is challenging you to build a model that predicts the total ride duration of taxi trips in New York City. Your primary dataset is one released by the NYC Taxi and Limousine Commission, which includes pickup time, geo-coordinates, number of passengers, and several other variables.

The competition dataset is based on the 2016 NYC Yellow Cab trip record data made available in Big Query on Google Cloud Platform. The data was originally published by the NYC Taxi and Limousine Commission (TLC). The data was sampled and cleaned for the purposes of this playground competition. Based on individual trip attributes, participants should predict the duration of each trip in the test set.

### File descriptions

- train.csv - the training set (contains 1458644 trip records)
- test.csv - the testing set (contains 625134 trip records)
- sample\_submission.csv - a sample submission file in the correct format  
<https://www.kaggle.com/c/6960/download-all>

### Data fields

Variable Name	Definition
id	a unique identifier for each trip
vendor_id	a code indicating the provider associated with the trip record
pickup_datetime	date and time when the meter was engaged
dropoff_datetime	date and time when the meter was disengaged
passenger_count	the number of passengers in the vehicle (driver entered value)
pickup_longitude	the longitude where the meter was engaged
pickup_latitude	the latitude where the meter was engaged
dropoff_longitude	the longitude where the meter was disengaged
dropoff_latitude	the latitude where the meter was disengaged
store_and_fwd_flag	This flag indicates whether the trip record was held in vehicle memory before sending to the vendor because the vehicle did not have a connection to the server: Y=store and forward; N=not a store and forward trip
trip_duration	duration of the trip in seconds

## Objective:

The purpose of this project is to build various models in an attempt to predict the trip duration of yellow taxis in New York City.

Using the techniques we've learned in the class, like classification, model diagnostics and transformation, we will explore data to find new patterns. And just like what is required in the Kaggle contest, we will try to predict the duration of each trip in the test set. We will build multiple linear regression modeling and then summary to interpret the results. We'll further analyze the results by adding discrimination in the model and then assess the discrimination with ROC curve.

vvv REWRITE vvv

# DATA EXPLORATION

## Load data

### View Data

```
## Observations: 1,458,644
## Variables: 11
## $ id          <chr> "id2875421", "id2377394", "id3858529", "id350467...
## $ vendor_id   <int> 2, 1, 2, 2, 2, 2, 1, 2, 1, 2, 2, 2, 2, 2, 1, ...
## $ pickup_datetime <chr> "2016-03-14 17:24:55", "2016-06-12 00:43:35", "2...
## $ dropoff_datetime <chr> "2016-03-14 17:32:30", "2016-06-12 00:54:38", "2...
## $ passenger_count <int> 1, 1, 1, 1, 1, 6, 4, 1, 1, 1, 1, 4, 2, 1, 1, 1, ...
## $ pickup_longitude <dbl> -73.98215, -73.98042, -73.97903, -74.01004, -73...
## $ pickup_latitude <dbl> 40.76794, 40.73856, 40.76394, 40.71997, 40.79321...
## $ dropoff_longitude <dbl> -73.96463, -73.99948, -74.00533, -74.01227, -73...
## $ dropoff_latitude <dbl> 40.76560, 40.73115, 40.71009, 40.70672, 40.78252...
## $ store_and_fwd_flag <chr> "N", "N", "N", "N", "N", "N", "N", "N", "N", "N"...
## $ trip_duration <int> 455, 663, 2124, 429, 435, 443, 341, 1551, 255, 1...
```

## Data Summary

```
## [1] "id"                "vendor_id"          "pickup_datetime"
## [4] "dropoff_datetime"   "passenger_count"    "pickup_longitude"
## [7] "pickup_latitude"    "dropoff_longitude"  "dropoff_latitude"
## [10] "store_and_fwd_flag" "trip_duration"
```

```
## [1] "id"                "vendor_id"          "pickup_datetime"
## [4] "passenger_count"   "pickup_longitude"   "pickup_latitude"
## [7] "dropoff_longitude" "dropoff_latitude"   "store_and_fwd_flag"
```

```
## Observations: 2,083,778
```

```
## Variables: 12
```

```
## $ id          <chr> "id2875421", "id2377394", "id3858529", "id350467...
## $ vendor_id    <int> 2, 1, 2, 2, 2, 2, 1, 2, 1, 2, 2, 2, 2, 2, 1, ...
## $ pickup_datetime <chr> "2016-03-14 17:24:55", "2016-06-12 00:43:35", "2...
## $ dropoff_datetime <chr> "2016-03-14 17:32:30", "2016-06-12 00:54:38", "2...
## $ passenger_count <int> 1, 1, 1, 1, 1, 6, 4, 1, 1, 1, 1, 4, 2, 1, 1, 1, ...
## $ pickup_longitude <dbl> -73.98215, -73.98042, -73.97903, -74.01004, -73....
## $ pickup_latitude <dbl> 40.76794, 40.73856, 40.76394, 40.71997, 40.79321...
## $ dropoff_longitude <dbl> -73.96463, -73.99948, -74.00533, -74.01227, -73....
## $ dropoff_latitude <dbl> 40.76560, 40.73115, 40.71009, 40.70672, 40.78252...
## $ store_and_fwd_flag <chr> "N", "N", "N", "N", "N", "N", "N", "N", "N"...
## $ trip_duration  <int> 455, 663, 2124, 429, 435, 443, 341, 1551, 255, 1...
## $ dset           <fct> train, train, train, train, train, train, train,...
```


```
##      id          vendor_id    pickup_datetime    dropoff_datetime
## Length:2083778    Min.      :1.000    Length:2083778    Length:2083778
## Class :character  1st Qu.:1.000    Class :character  Class :character
## Mode  :character  Median :2.000    Mode  :character  Mode  :character
##
##              Mean      :1.535
##              3rd Qu.:2.000
##              Max.     :2.000
##
## passenger_count pickup_longitude pickup_latitude dropoff_longitude
## Min.      :0.000    Min.      :-121.93    Min.      :34.36    Min.      :-121.93
## 1st Qu.:1.000    1st Qu.: -73.99    1st Qu.:40.74    1st Qu.: -73.99
## Median :1.000    Median : -73.98    Median :40.75    Median : -73.98
## Mean      :1.664    Mean      : -73.97    Mean      :40.75    Mean      : -73.97
## 3rd Qu.:2.000    3rd Qu.: -73.97    3rd Qu.:40.77    3rd Qu.: -73.96
## Max.      :9.000    Max.       : -61.34    Max.      :51.88    Max.       : -61.34
##
## dropoff_latitude store_and_fwd_flag trip_duration      dset
## Min.      :32.18    Length:2083778    Min.      :      1    test : 625134
## 1st Qu.:40.74    Class :character  1st Qu.:    397    train:1458644
## Median :40.75    Mode  :character  Median :    662
## Mean      :40.75                                Mean      :    959
## 3rd Qu.:40.77                                3rd Qu.:   1075
## Max.      :48.86                                Max.      :3526282
##
##                                NA's      :625134
```

## Missing values

```
##      columns na_count neg_count zero_count unique_count
## 1          id         0         0         NA       2083778
```

```
## 2      vendor_id      0      0      0      2
## 3  pickup_datetime      0      0     NA    1926217
## 4  dropoff_datetime  625134     NA     NA    1380378
## 5    passenger_count      0      0     83      10
## 6  pickup_longitude      0  2083778      0    24960
## 7  pickup_latitude      0      0      0    48068
## 8  dropoff_longitude      0  2083778      0    36977
## 9  dropoff_latitude      0      0      0    67086
## 10 store_and_fwd_flag      0      0     NA      2
## 11    trip_duration  625134     NA     NA    7418
## 12          dset      0     NA      0      2
```

```
##          columns na_count neg_count zero_count unique_count
## 1            id      0      0      NA    2083778
## 2      vendor_id      0      0      0      2
## 3  pickup_datetime      0      0     NA    1926217
## 4  dropoff_datetime  625134     NA     NA    1380378
## 5    passenger_count      0      0     83      10
## 6  pickup_longitude      0  2083778      0    24960
## 7  pickup_latitude      0      0      0    48068
## 8  dropoff_longitude      0  2083778      0    36977
## 9  dropoff_latitude      0      0      0    67086
## 10 store_and_fwd_flag      0      0     NA      2
## 11    trip_duration  625134     NA     NA    7418
## 12          dset      0     NA      0      2
```



```
##      Mode  FALSE  TRUE
## logical 625134 1458644
```

#Data Preparation

## Reformatting features

For our following analysis, we will turn the data and time from characters into *date* objects. We also recode *vendor\_id* as a factor. This makes it easier to visualise relationships that involve these features.

## Visualizations




Fig. 2

```
## train$store_and_fwd_flag: N
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.000  5.984   6.495   6.464  6.979  15.076
## -----
## train$store_and_fwd_flag: Y
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.6931  6.1203  6.6995  6.6274  7.2442  9.2087

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

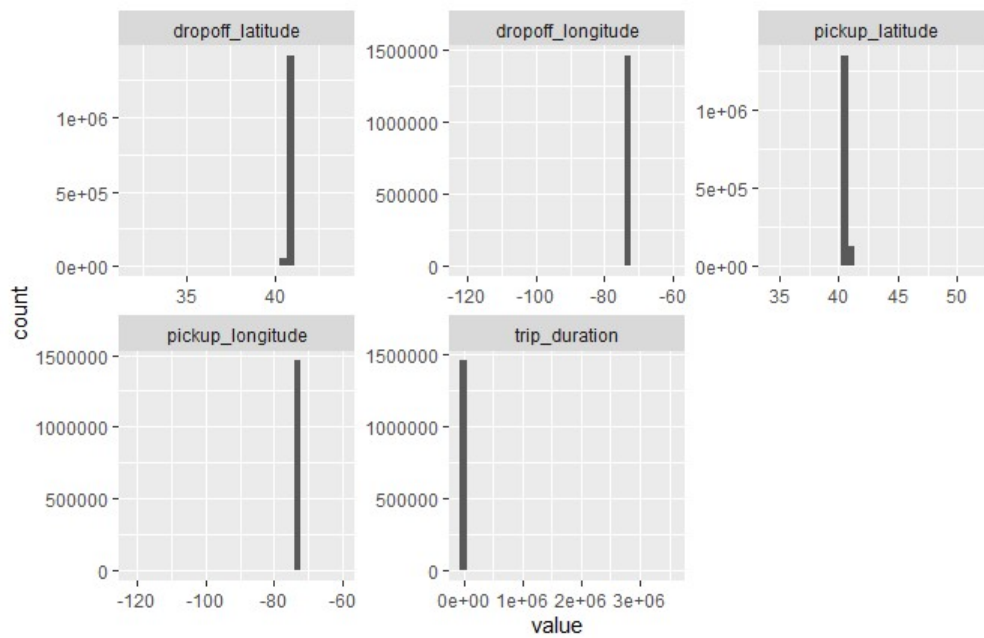
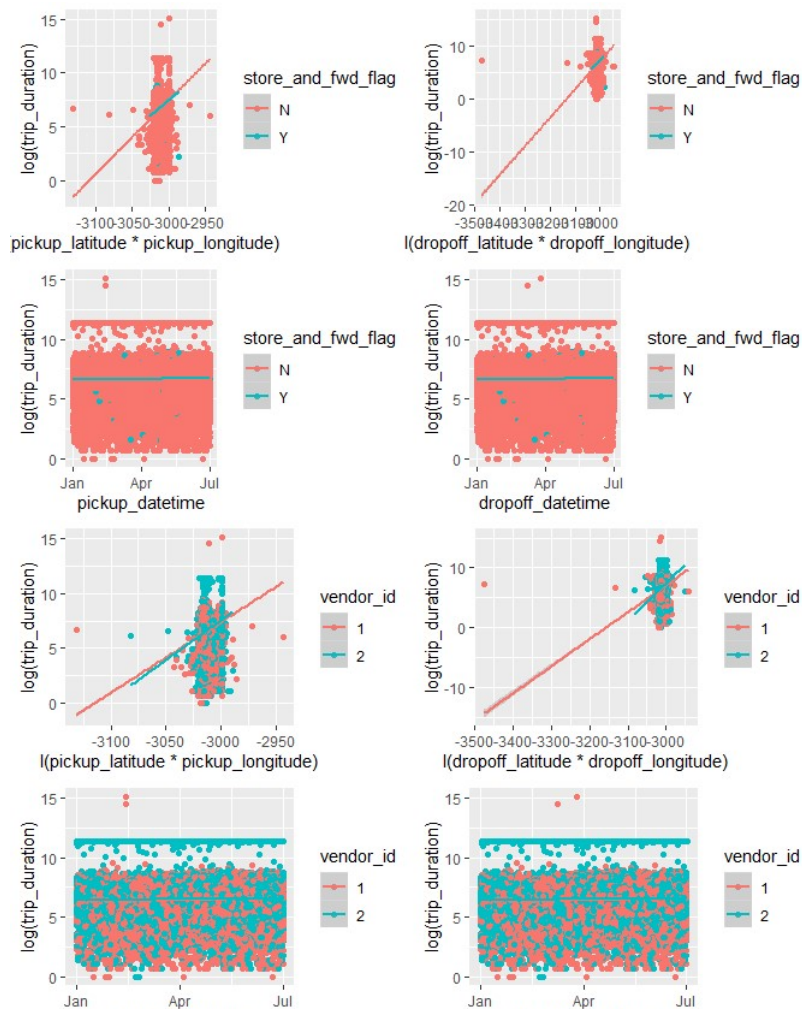
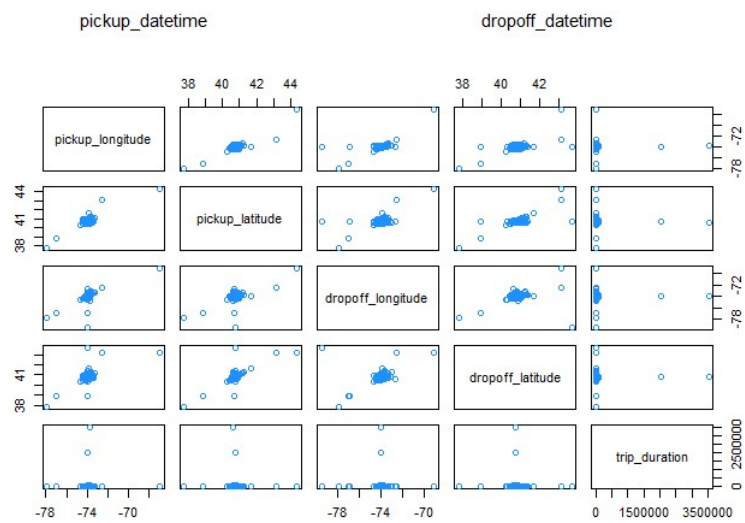


Fig. 2

## Trip duration vs pickup and dropoff datetime and location using 20% of sampled data





## More visualizations

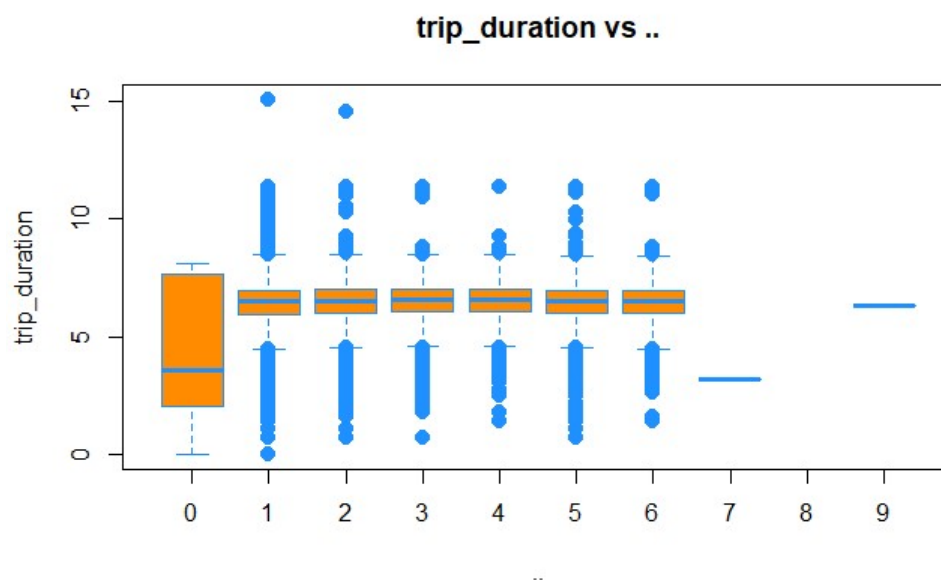


Fig. 2

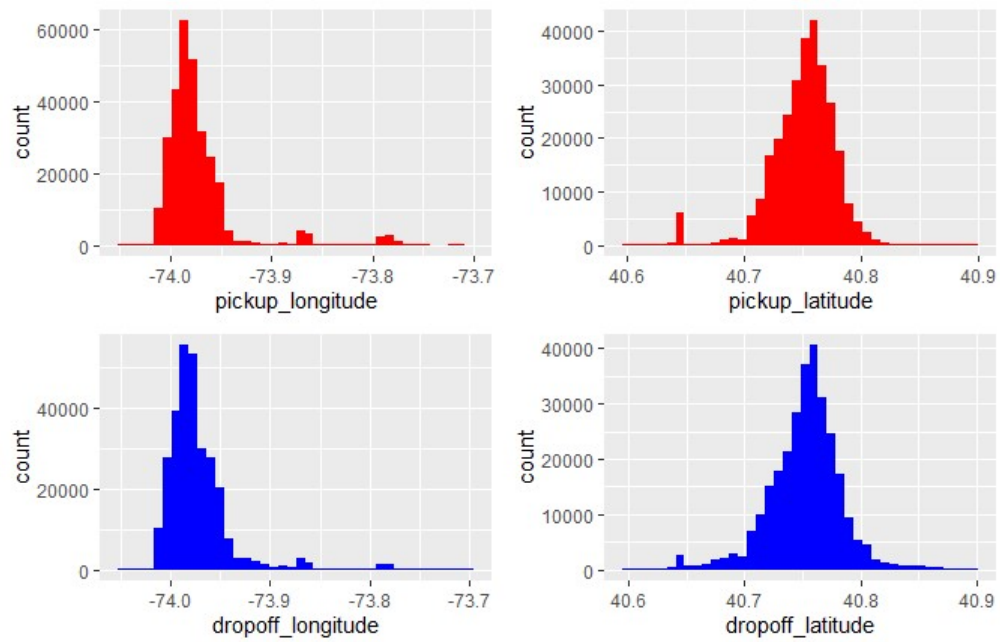


Fig. 6




Fig. 30a



## An excursion into classification

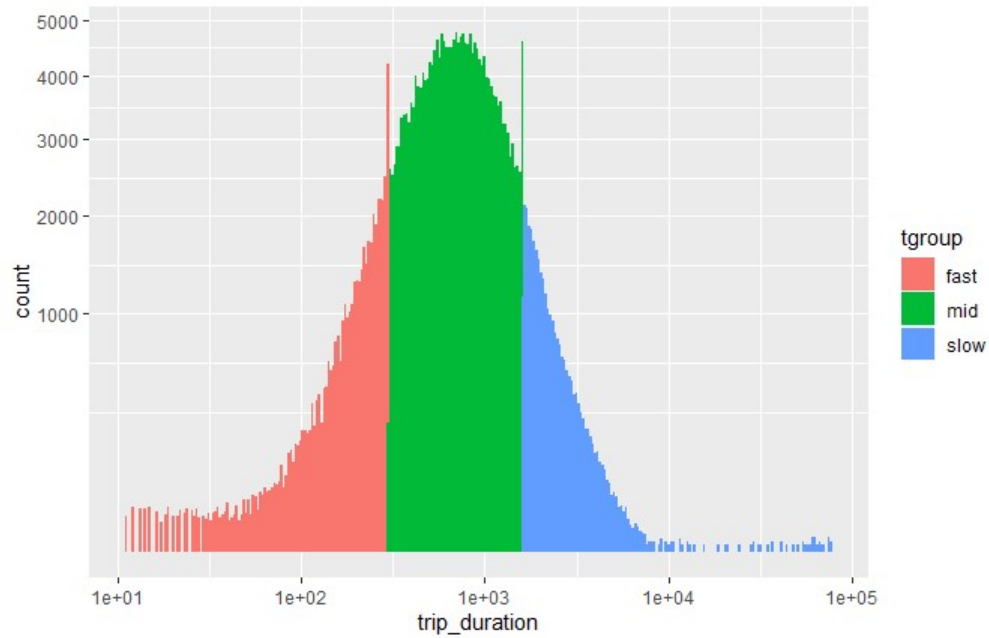


Fig. 31




Fig. 32




Fig. 34

## Model, correlation




Fig. 35




Fig. 36

## Modeling

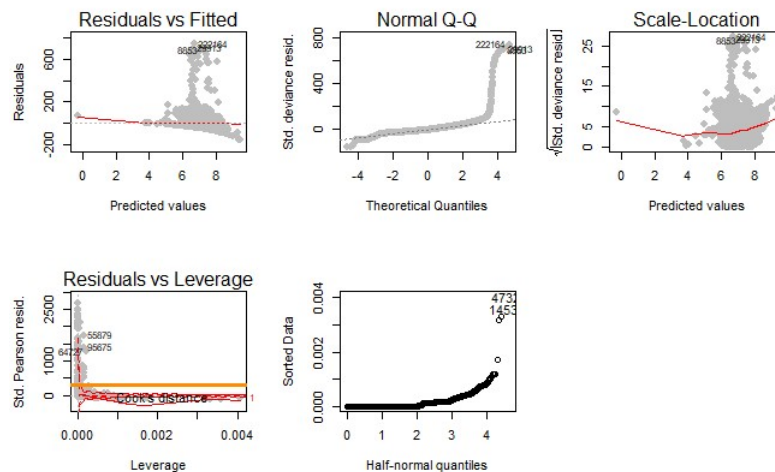
### Sample Model 1 Poisson (No store\_and\_fwd\_flag, id or vendor\_id)

```
##
## Call:
## glm(formula = trip_duration ~ pickup_datetime:dropoff_datetime +
##      pickup_latitude:pickup_longitude + dropoff_latitude:dropoff_longitude,
##      family = poisson)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -156.81  -15.43   -4.92    8.35   740.39
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      3.149e+02  6.351e-02 4957.2  <2e-16 ***
## pickup_datetime:dropoff_datetime  3.183e-18  4.902e-21  649.3  <2e-16 ***
## pickup_latitude:pickup_longitude  6.069e-02  1.666e-05 3643.4  <2e-16 ***
## dropoff_latitude:dropoff_longitude 4.378e-02  2.015e-05 2172.3  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 136608788  on 290041  degrees of freedom
## Residual deviance: 114317748  on 290038  degrees of freedom
## AIC: 116727043
##
## Number of Fisher Scoring iterations: 5

## # A tibble: 4 x 5
##   term                                estimate std.error statistic p.value
##   <chr>                                <dbl>     <dbl>     <dbl>  <dbl>
```

```
## 1 (Intercept) 3.15e+ 2 6.35e- 2 4957. 0
## 2 pickup_datetime:dropoff_datetime 3.18e-18 4.90e-21 649. 0
## 3 pickup_latitude:pickup_longitude 6.07e- 2 1.67e- 5 3643. 0
## 4 dropoff_latitude:dropoff_longitude 4.38e- 2 2.02e- 5 2172. 0

## # A tibble: 1 x 7
## null.deviance df.null logLik AIC BIC deviance df.residual
## <dbl> <int> <dbl> <dbl> <dbl> <dbl> <int>
## 1 136608788. 290041 -58363517. 116727043. 116727085. 114317748. 290038
```




Breusch-Pagan Test.


Cooks Distance H0: Homoscedasticity. The errors have constant variance about the true model. H1: Heteroscedasticity. The errors have non-constant variance about the true model. Leverage, Outliers, Influence, coef Change

```
## 11
## TRUE

##
## studentized Breusch-Pagan test
##
## data: sample_model
## BP = 123.87, df = 3, p-value < 2.2e-16


## [1] 257881
```





```
##          (Intercept)  pickup_datetime:dropoff_datetime
##          3.148518e+02          3.182603e-18
## pickup_latitude:pickup_longitude dropoff_latitude:dropoff_longitude
##          6.069042e-02          4.377879e-02
```

```
##          (Intercept) dropoff_longitude
##          182176.241          2449.298
```



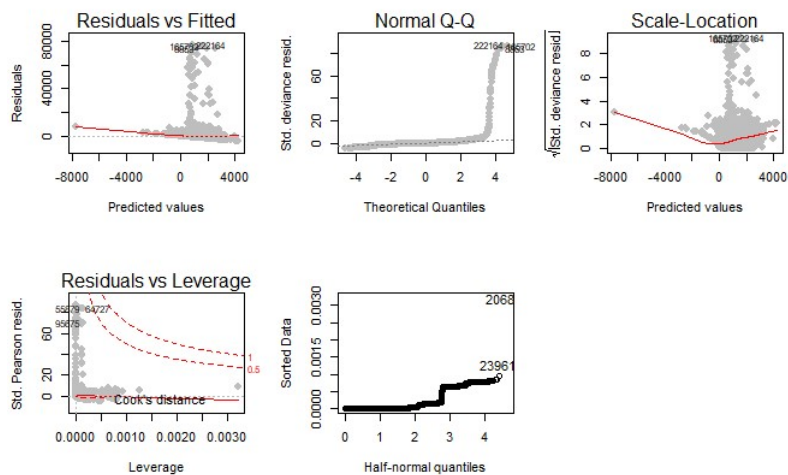
## Sample Model 2 Gaussian(No id or vendor\_id)

```
##
## Call:
## glm(formula = trip_duration ~ pickup_datetime:dropoff_datetime +
##      pickup_latitude:pickup_longitude + dropoff_latitude:dropoff_longitude +
##      store_and_fwd_flag, family = gaussian)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
##    -4148    -392    -127     251    76335
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    3.795e+05  2.263e+03  167.699 < 2e-16 ***
## store_and_fwd_flagY    1.810e+02  2.231e+01   8.114 4.93e-16 ***
## pickup_datetime:dropoff_datetime    2.710e-15  1.257e-16  21.560 < 2e-16 ***
## pickup_latitude:pickup_longitude    8.126e+01  6.250e-01  130.025 < 2e-16 ***
## dropoff_latitude:dropoff_longitude    4.628e+01  6.545e-01  70.702 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## (Dispersion parameter for gaussian family taken to be 774496.8)
##
## Null deviance: 2.4925e+11 on 290041 degrees of freedom
## Residual deviance: 2.2463e+11 on 290037 degrees of freedom
## AIC: 4756071
##
## Number of Fisher Scoring iterations: 2
```

```
## # A tibble: 5 x 5
##   term                                estimate std.error statistic  p.value
##   <chr>                                <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)                        3.80e+ 5  2.26e+ 3    168.     0.
## 2 store_and_fwd_flagY                 1.81e+ 2  2.23e+ 1     8.11  4.93e- 16
## 3 pickup_datetime:dropoff_datetime    2.71e-15  1.26e-16    21.6   5.16e-103
## 4 pickup_latitude:pickup_longitude     8.13e+ 1  6.25e- 1    130.     0.
## 5 dropoff_latitude:dropoff_longitude   4.63e+ 1  6.55e- 1    70.7     0.
```

```
## # A tibble: 1 x 7
##   null.deviance df.null logLik AIC BIC deviance df.residual
##   <dbl> <int> <dbl> <dbl> <dbl> <dbl> <int>
## 1 249251879377. 290041 -2378030. 4756071. 4756134. 224632723577. 290037
```



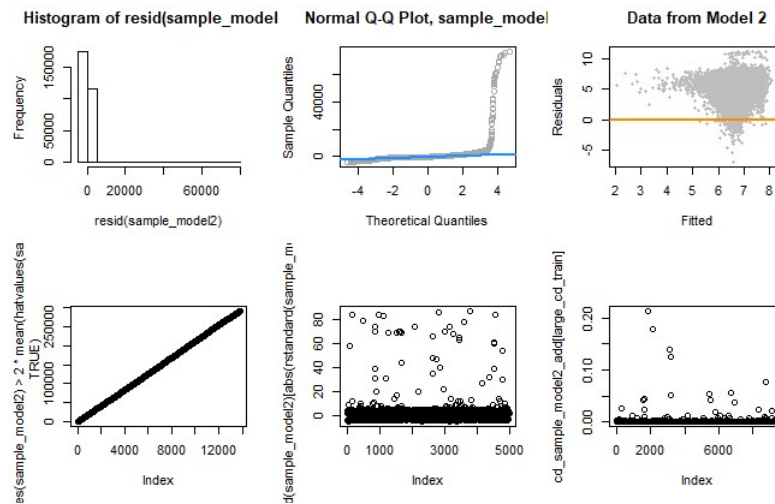
Breusch-Pagan Test.

Cooks Distance H0: Homoscedasticity. The errors have constant variance about the true model. H1: Heteroscedasticity. The errors have non-constant variance about the true model. Leverage, Outliers, Influence, coef Change

```
## 11
## FALSE
```


```
##
## studentized Breusch-Pagan test
##
## data: sample_model2
## BP = 123.86, df = 4, p-value < 2.2e-16
```

```
## [1] 9498
```



```
##               (Intercept)               store_and_fwd_flagY
##      3.795352e+05                1.810469e+02
## pickup_datetime:dropoff_datetime pickup_latitude:pickup_longitude
##      2.709699e-15                8.126130e+01
## dropoff_latitude:dropoff_longitude
##      4.627524e+01
```

```
##      (Intercept) dropoff_longitude
##      79730.086      1064.841
```



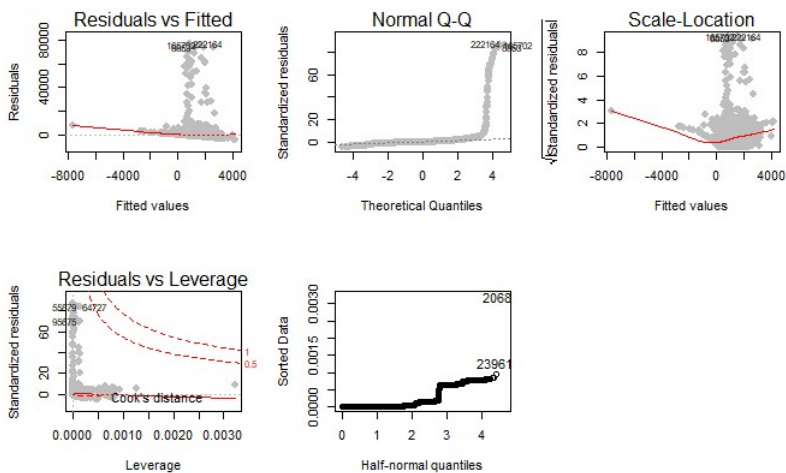
## Sample Model 3 Negative Binomial (No id)

```
##
## Call:
## lm(formula = trip_duration ~ pickup_datetime:dropoff_datetime +
##      pickup_latitude:pickup_longitude + dropoff_latitude:dropoff_longitude +
##      store_and_fwd_flag + vendor_id, family = negative.binomial(1))
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4133    -392    -127     251   76323
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   3.794e+05  2.263e+03  167.672 < 2e-16 ***
## store_and_fwd_flagY 1.952e+02  2.238e+01   8.722 < 2e-16 ***
## vendor_id2      2.629e+01  3.286e+00   7.998 1.27e-15 ***
## pickup_datetime:dropoff_datetime 2.717e-15  1.257e-16  21.619 < 2e-16 ***
## pickup_latitude:pickup_longitude 8.122e+01  6.249e-01 129.971 < 2e-16 ***
## dropoff_latitude:dropoff_longitude 4.629e+01  6.544e-01  70.735 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 880 on 290036 degrees of freedom
## Multiple R-squared:  0.09897,    Adjusted R-squared:  0.09896
## F-statistic: 6372 on 5 and 290036 DF,  p-value: < 2.2e-16
```

```
## # A tibble: 6 x 5
##   term                estimate std.error statistic  p.value
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)        3.79e+ 5  2.26e+ 3    168.    0.
## 2 store_and_fwd_flagY 1.95e+ 2  2.24e+ 1     8.72 2.75e- 18
## 3 vendor_id2         2.63e+ 1  3.29e+ 0     8.00 1.27e- 15
## 4 pickup_datetime:dropoff_datetime 2.72e-15  1.26e-16    21.6 1.45e-103
## 5 pickup_latitude:pickup_longitude 8.12e+ 1  6.25e- 1    130.    0.
## 6 dropoff_latitude:dropoff_longitude 4.63e+ 1  6.54e- 1    70.7    0.
```

```
## # A tibble: 1 x 11
##   r.squared adj.r.squared sigma statistic p.value    df logLik   AIC   BIC
##   <dbl>      <dbl>    <dbl>    <dbl>    <dbl> <int>  <dbl> <dbl> <dbl>
## 1  0.0990      0.0990    880.    6372.    0      6 -2.38e6 4.76e6 4.76e6
## # ... with 2 more variables: deviance <dbl>, df.residual <int>
```



Breusch-Pagan Test.

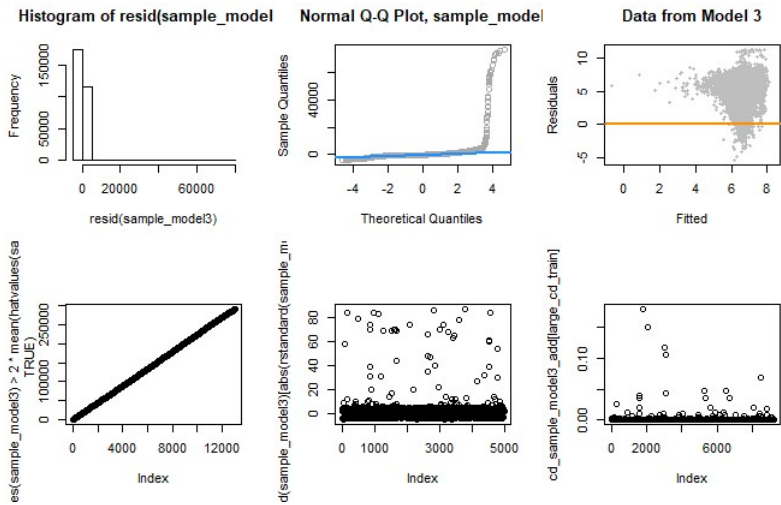


Cooks Distance H0: Homoscedasticity. The errors have constant variance about the true model. H1: Heteroscedasticity. The errors have non-constant variance about the true model. Leverage, Outliers, Influence, coef Change

```
## 11
## FALSE


##
## studentized Breusch-Pagan test
##
## data: sample_model3
## BP = 150.33, df = 5, p-value < 2.2e-16


## [1] 9185
```



```
## (Intercept) store_and_fwd_flagY
## 3.794375e+05 1.952059e+02
## vendor_id2 pickup_datetime:dropoff_datetime
## 2.628544e+01 2.716875e-15
## pickup_latitude:pickup_longitude dropoff_latitude:dropoff_longitude
## 8.122143e+01 4.629244e+01

## (Intercept) dropoff_longitude
## 79811.38 1065.94
```






	Model 1	Model 2	Model 3
AIC	116727042.65499	4756071.0066072	4756009.03978159
BIC	116727084.966114	4756134.47329331	4756083.08424872

## CONCLUSION

With 3 models computed, we select the model with the lowest combination of AIC and BIC. From the table, we can see the model to pick is model 1

Model 1 showed the best result. We can observe its performance by plotting the datasets Vendor\_ID values against the predicted values.



## APPENDIX

### Code used in analysis

```
#List.of.packages <-
  c("alluvial", "caret", "caret", "corrplot", "corrplot", "data.table", "dplyr", "faraway", "forcats", "geosphere", "gg")
#new.packages <- list.of.packages[!(list.of.packages %in% installed.packages())
  [, "Package"]]
#if(length(new.packages)) install.packages(new.packages)
require(knitr)
knitr::opts_chunk$set(echo = FALSE, warning = FALSE, fig.align='center')

library(faraway)
library(MASS)
library(psych)
library(pROC)
library(corrplot)
library(jtools)
library(mice)
library('corr')
```

```

library(kableExtra)
library(gridExtra)
library(pander)
library(zoo)
library(lmtest)
library(corr)
library(broom)

#devtools::install_github("thomasp85/patchwork")
library(patchwork)
library(tidyverse)
library(ggplot2)
library(ggplot2)
library(reshape2)
library(leaps)
library(caret)
library(naniar)
library('ggplot2') # visualisation
library('scales') # visualisation
library('grid') # visualisation
library('RColorBrewer') # visualisation
library('corrplot') # visualisation
library('alluvial') # visualisation
library('dplyr') # data manipulation
library('readr') # input/output
library('data.table') # data manipulation
library('tibble') # data wrangling
library('tidyr') # data wrangling
library('stringr') # string manipulation
library('forcats') # factor manipulation
library('lubridate') # date and time
library('geosphere') # geospatial locations
library('leaflet') # maps
library('leaflet.extras') # maps
library('maps') # maps
library('xgboost') # modelling
library('caret') # modelling
library('widgetframe') #visualizaiton
library('grid')
library('gridExtra')
# Define multiple plot function
#
# ggplot objects can be passed in ..., or to plotlist (as a list of ggplot objects)
# - cols:   Number of columns in layout
# - layout: A matrix specifying the layout. If present, 'cols' is ignored.
#
# If the layout is something like matrix(c(1,2,3,3), nrow=2, byrow=TRUE),
# then plot 1 will go in the upper left, 2 will go in the upper right, and
# 3 will go all the way across the bottom.
#

multiplot <- function(..., plotlist=NULL, file, cols=1, layout=NULL) {

  # Make a list from the ... arguments and plotlist
  plots <- c(list(...), plotlist)

  numPlots = length(plots)

  # If layout is NULL, then use 'cols' to determine layout
  if (is.null(layout)) {
    # Make the panel
    # ncol: Number of columns of plots
    # nrow: Number of rows needed, calculated from # of cols

```

```

    layout <- matrix(seq(1, cols * ceiling(numPlots/cols)),
                     ncol = cols, nrow = ceiling(numPlots/cols))
  }

  if (numPlots==1) {
    print(plots[[1]])
  } else {
    # Set up the page
    grid.newpage()
    pushViewport(viewport(layout = grid.layout(nrow(layout), ncol(layout))))

    # Make each plot, in the correct location
    for (i in 1:numPlots) {
      # Get the i,j matrix positions of the regions that contain this subplot
      matchidx <- as.data.frame(which(layout == i, arr.ind = TRUE))

      print(plots[[i]], vp = viewport(layout.pos.row = matchidx$row,
                                       layout.pos.col = matchidx$col))
    }
  }
}

var_stats<- function(df){
  wt<-data.frame(columns=colnames(df))
  wt$na_count <- sapply(df, function(y) sum(is.na(y)))
  wt$neg_count <- sapply(df, function(y) sum(y<0))
  wt$zero_count <- sapply(df, function(y) sum(as.integer(y)==0))
  wt$unique_count <- sapply(df, function(y) sum(n_distinct(y)))
  print(wt)
  return(wt)
}

rows <-
  c("id", "vendor_id", "pickup_datetime", "dropoff_datetime", "passenger_count", "pickup_longitude", "pickup_latitude",
    "dropoff_longitude", "dropoff_latitude", "store_and_fwd_flag", "trip_duration")

def <- c("a unique identifier for each trip",
"a code indicating the provider associated with the trip record",
"date and time when the meter was engaged",
"date and time when the meter was disengaged",
"the number of passengers in the vehicle (driver entered value)",
"the longitude where the meter was engaged",
"the latitude where the meter was engaged",
"the longitude where the meter was disengaged",
"the latitude where the meter was disengaged",
"This flag indicates whether the trip record was held in vehicle memory before sending to
the vendor because the vehicle did not have a connection to the server: Y=store and
forward; N=not a store and forward trip",
"duration of the trip in seconds")

kable(cbind(rows, def), col.names = c("Variable Name", "Definition")) %>% kable_styling()
train <- as_tibble(fread('data/train.csv'))
test <- as_tibble(fread('data/test.csv'))
sample_submit <- as_tibble(fread('data/sample_submission.csv'))
#str(train)
#glimpse(train)
#summary(train)
#describe(train)
names(train)
names(test)
#glimpse(test)

#
vars_to_add <- train[!names(train) %in% names(test)]

```

```

#vvvvv
## Combining train and test

combine <- rbind(train %>% mutate(dset = "train"),
                 test %>% mutate(dset = "test",
                                dropoff_datetime = NA,
                                trip_duration = NA))

combine <- combine %>% mutate(dset = factor(dset))
glimpse(combine)
summary(combine)
var_stats(combine)
gg_miss_upset(combine)
summary(complete.cases(combine))
train <- train %>%
  mutate(pickup_datetime = ymd_hms(pickup_datetime),
         dropoff_datetime = ymd_hms(dropoff_datetime),
         vendor_id = factor(vendor_id),
         passenger_count = factor(passenger_count))

#ggplot(combine, aes(trip_duration)) +
#  geom_histogram(aes(y = ..density..))

attach(train)
boxplot(by(log(train$trip_duration), train$store_and_fwd_flag, summary), col=c("red", "blue"), xlab="Store
and Forward Flag", ylab="Trip Duration")
by(log(train$trip_duration), train$store_and_fwd_flag, summary)

#plot(trip_duration ~ dropoff_longitude, pch = 20, cex = 2, col = "grey")

train[sapply(train, function(x) is.numeric(x) && !is.na(x))] %>%
  gather() %>%
  ggplot(aes(value), main="") +
  facet_wrap(~ key, scales = "free") +
  geom_histogram()

sub_train = train %>% sample_frac(.2)
attach(sub_train)
g1<-ggplot(sub_train, aes(x=I(pickup_latitude*pickup_longitude), y=log(trip_duration),
  color = store_and_fwd_flag)) + geom_point() + stat_smooth(method="glm", se=TRUE)
g2<-ggplot(sub_train, aes(x=I(dropoff_latitude*dropoff_longitude), y=log(trip_duration),
  color = store_and_fwd_flag)) + geom_point() + stat_smooth(method="glm", se=TRUE)
g3<-ggplot(sub_train, aes(x=pickup_datetime, y=log(trip_duration), color =
  store_and_fwd_flag)) + geom_point() + stat_smooth(method="glm", se=TRUE)
g4<-ggplot(sub_train, aes(x=dropoff_datetime, y=log(trip_duration), color =
  store_and_fwd_flag)) + geom_point() + stat_smooth(method="glm", se=TRUE)
grid.arrange(g1, g2, g3, g4, ncol = 2)

g1<-ggplot(sub_train, aes(x=I(pickup_latitude*pickup_longitude), y=log(trip_duration),
  color = vendor_id)) + geom_point() + stat_smooth(method="glm", se=TRUE)
g2<-ggplot(sub_train, aes(x=I(dropoff_latitude*dropoff_longitude), y=log(trip_duration),
  color = vendor_id)) + geom_point() + stat_smooth(method="glm", se=TRUE)
g3<-ggplot(sub_train, aes(x=pickup_datetime, y=log(trip_duration), color = vendor_id))
+ geom_point() + stat_smooth(method="glm", se=TRUE)
g4<-ggplot(sub_train, aes(x=dropoff_datetime, y=log(trip_duration), color = vendor_id))
+ geom_point() + stat_smooth(method="glm", se=TRUE)
grid.arrange(g1, g2, g3, g4, ncol = 2)

pairs(sub_train[sapply(sub_train, function(x) is.numeric(x))], col = "dodgerblue")

ssub_train<-sub_train[sapply(sub_train, function(x) is.numeric(x) && !is.na(x))]

ssub_train %>%
  correlate() %>%
  network_plot(min_cor = .2)

```

```

#log(sub_train$trip_duration) %>% as.double() %>% boxplot()
#bins
#scale_x_log10() +
#scale_y_sqrt()
attach(sub_train)
boxplot(log(trip_duration) ~ as.factor(passenger_count),
        xlab = "..",
        ylab = "trip_duration",
        main = "trip_duration vs ..",
        pch = 20,
        cex = 2,
        col = "darkorange",
        border = "dodgerblue")

p1 <- sub_train %>%
  filter(pickup_longitude > -74.05 & pickup_longitude < -73.7) %>%
  ggplot(aes(pickup_longitude)) +
  geom_histogram(fill = "red", bins = 40)

p2 <- sub_train %>%
  filter(dropoff_longitude > -74.05 & dropoff_longitude < -73.7) %>%
  ggplot(aes(dropoff_longitude)) +
  geom_histogram(fill = "blue", bins = 40)

p3 <- sub_train %>%
  filter(pickup_latitude > 40.6 & pickup_latitude < 40.9) %>%
  ggplot(aes(pickup_latitude)) +
  geom_histogram(fill = "red", bins = 40)

p4 <- sub_train %>%
  filter(dropoff_latitude > 40.6 & dropoff_latitude < 40.9) %>%
  ggplot(aes(dropoff_latitude)) +
  geom_histogram(fill = "blue", bins = 40)

layout <- matrix(c(1,2,3,4),2,2,byrow=FALSE)
multiplot(p1, p2, p3, p4, layout=layout)
p1 <- 1; p2 <- 1; p3 <- 1; p4 <- 1
#jfk_coord <- tibble(Lon = -73.778889, Lat = 40.639722)

#la_guardia_coord <- tibble(Lon = -73.872611, Lat = 40.77725)

#train$jfk_dist_pick <- distCosine(pick_coord, jfk_coord)

#train$jfk_dist_drop <- distCosine(drop_coord, jfk_coord)

#train$lg_dist_pick <- distCosine(pick_coord, la_guardia_coord)

#train$lg_dist_drop <- distCosine(drop_coord, la_guardia_coord)

pick_coord <- sub_train %>% select(pickup_longitude, pickup_latitude)

drop_coord <- sub_train %>% select(dropoff_longitude, dropoff_latitude)

sub_train$dist <- distCosine(pick_coord, drop_coord)

#train$bearing = bearing(pick_coord, drop_coord)

```

```

sub_train <- sub_train %>%

  mutate(speed = dist/trip_duration*3.6,

         date = date(pickup_datetime),

         month = month(pickup_datetime, label = TRUE),

         wday = wday(pickup_datetime, label = TRUE),

         wday = fct_relevel(wday, c("Mon", "Tue", "Wed", "Thu", "Fri", "Sat", "Sun")),

         hour = hour(pickup_datetime),

         work = (hour %in% seq(8,18)) & (wday %in% c("Mon","Tue","Wed","Thu","Fri")),

#       jfk_trip = (jfk_dist_pick < 2e3) | (jfk_dist_drop < 2e3),

#       lg_trip = (lg_dist_pick < 2e3) | (lg_dist_drop < 2e3),

#       blizzard = !( (date < ymd("2016-01-22") | (date > ymd("2016-01-29"))) )

  )

```

```

sub_train <- sub_train %>%

  filter(trip_duration < 22*3600,

         dist > 0 | (near(dist, 0) & trip_duration < 60),

#       jfk_dist_pick < 3e5 & jfk_dist_drop < 3e5,

         trip_duration > 10,

         speed < 100)

```

```

sub_train %>%

  select(-id, -pickup_datetime, -dropoff_datetime, -date) %>% #-jfk_dist_pick,

#       -jfk_dist_drop, -lg_dist_pick, -lg_dist_drop, -date) %>%

  mutate(passenger_count = as.integer(passenger_count),

         vendor_id = as.integer(vendor_id),

         store_and_fwd_flag = as.integer(as.factor(store_and_fwd_flag)),

#       jfk_trip = as.integer(jfk_trip),

         wday = as.integer(wday),

         month = as.integer(month),

         work = as.integer(work))%>%

#       lg_trip = as.integer(lg_trip),

```

```

#       blizzard = as.integer(blizzard),

#       has_snow = as.integer(has_snow),

#       has_rain = as.integer(has_rain)) %>%
#
#   select(trip_duration, speed, everything()) %>%

#   cor(use="complete.obs", method = "spearman") %>%

#   corplot(type="lower", method="circle", diag=FALSE)

train_group <- sub_train %>%

  mutate(tgroup = case_when(trip_duration < 3e2 ~ "fast",

                             trip_duration >= 3e2 & trip_duration <= 1.6e3 ~ "mid",

                             trip_duration > 1.6e3 ~ "slow"))

train_group %>%

  ggplot(aes(trip_duration, fill = tgroup)) +

  geom_histogram(bins = 300) +

  scale_x_log10() +

  scale_y_sqrt()

train_group <- train_group %>%

  filter(tgroup != "mid")

p1 <- train_group %>%

  ggplot(aes(wday, fill = tgroup)) +

  geom_bar(position = "fill") +

  theme(legend.position = "none")

p2 <- train_group %>%

  ggplot(aes(month, fill = tgroup)) +

  geom_bar(position = "fill") +

  theme(legend.position = "none")

```



```

p3 <- train_group %>%

  ggplot(aes(hour, fill = tgroup)) +

  geom_bar(position = "fill")

p7 <- train_group %>%

  ggplot(aes(work, fill = tgroup)) +

  geom_bar(position = "fill") +

  theme(legend.position = "none")


layout <- matrix(c(1,1,2,2,3,3,3,3,4,5,6,7),3,4,byrow=TRUE)

multiplot(p1, p2, p7, layout=layout)

p1 <- 1; p2 <- 1; p7 <- 1


allu_train <- train_group %>%

  group_by(tgroup, work, wday) %>% # jfk_trip

  count() %>%

  ungroup


alluvial(allu_train %>% select(-n),

  freq=allu_train$n, border=NA,

  col=ifelse(allu_train$tgroup == "fast", "red", "blue"),

  cex=0.75,

  hide = allu_train$n < 150,

  ordering = list(

    order(allu_train$tgroup=="fast"),

    # NULL,

    NULL,

    NULL))


foo <- combine %>%

  mutate(date = date(ymd_hms(pickup_datetime))) %>%

  group_by(date, dset) %>%

  count() %>%

```

```

ungroup()

foo %>%

  ggplot(aes(date, n/1e3, color = dset)) +

  geom_line(size = 1.5) +

  labs(x = "", y = "Kilo trips per day")

pick_good <- combine %>%

  filter(pickup_longitude > -75 & pickup_longitude < -73) %>%

  filter(pickup_latitude > 40 & pickup_latitude < 42)

pick_good <- sample_n(pick_good, 5e3)

pick_good %>%

  ggplot(aes(pickup_longitude, pickup_latitude, color = dset)) +

  geom_point(size=0.1, alpha = 0.5) +

  coord_cartesian(xlim = c(-74.02,-73.77), ylim = c(40.63,40.84)) +

  facet_wrap(~ dset) +

  #guides(color = guide_legend(override.aes = list(alpha = 1, size = 4))) +

  theme(legend.position = "none")

attach(sub_train)
sample_model = glm(trip_duration ~
  pickup_datetime:dropoff_datetime+pickup_latitude:pickup_longitude+dropoff_latitude:dropoff_longitude,
  family = poisson)

par(mfrow = c(2,3))
plot(sample_model,
  pch = 20,
  cex = 2,
  col = "grey")
abline(sample_model, lwd = 3, col = "darkorange")
halfnorm(hatvalues(sample_model))

summary(sample_model)
#confint(sample_model, Level = 0.99)
tidy(sample_model)
#augment(sample_model)
glance(sample_model)
cooks.distance(sample_model)[11] > 4 / length(cooks.distance(sample_model))
bptest(sample_model)
par(mfrow = c(2,3))
hist(resid(sample_model))

qqnorm(resid(sample_model), main = "Normal Q-Q Plot, sample_model", col = "darkgrey")
qqline(resid(sample_model), col = "dodgerblue", lwd = 2)

```

```

plot(log(fitted(sample_model)), log(resid(sample_model)), col = "grey", pch = 20,
      xlab = "Fitted", ylab = "Residuals", main = "Data from Model 1")
abline(h = 0, col = "darkorange", lwd = 2)

plot(which(hatvalues(sample_model) > 2 * mean(hatvalues(sample_model))), TRUE))

plot(rstandard(sample_model)[abs(rstandard(sample_model)) > 2])
cd_sample_model_add = cooks.distance(sample_model)
sum(cd_sample_model_add > 4 / length(cd_sample_model_add))

large_cd_train = cd_sample_model_add > 4 / length(cd_sample_model_add)
plot(cd_sample_model_add[large_cd_train])

coef(sample_model)
sample_model_add_fix = lm(trip_duration ~ dropoff_longitude,
                          data = train,
                          subset = cd_sample_model_add <= 4 / length(cd_sample_model_add))
coef(sample_model_add_fix)

#set.seed(42)
#shapiro.test(resid(sample_model))

boxcox(sample_model, plotit = TRUE)
attach(sub_train)
sample_model2 = glm(trip_duration ~
  pickup_datetime:dropoff_datetime+pickup_latitude:pickup_longitude+dropoff_latitude:dropoff_longitude+store_ai
  family = gaussian)

par(mfrow = c(2,3))
plot(sample_model2,
      pch = 20,
      cex = 2,
      col = "grey")
abline(sample_model2, lwd = 3, col = "darkorange")
halfnorm(hatvalues(sample_model2))

summary(sample_model2)
#confint(sample_model2, Level = 0.99)
tidy(sample_model2)
#augment(sample_model)
glance(sample_model2)
cooks.distance(sample_model2)[11] > 4 / length(cooks.distance(sample_model2))
bptest(sample_model2)
par(mfrow = c(2,3))
hist(resid(sample_model2))

qqnorm(resid(sample_model2), main = "Normal Q-Q Plot, sample_model", col = "darkgrey")
qqline(resid(sample_model2), col = "dodgerblue", lwd = 2)

plot(log(fitted(sample_model2)), log(resid(sample_model2)), col = "grey", pch = 20,
      xlab = "Fitted", ylab = "Residuals", main = "Data from Model 2")
abline(h = 0, col = "darkorange", lwd = 2)

plot(which(hatvalues(sample_model2) > 2 * mean(hatvalues(sample_model2))), TRUE))

plot(rstandard(sample_model2)[abs(rstandard(sample_model2)) > 2])
cd_sample_model2_add = cooks.distance(sample_model2)
sum(cd_sample_model2_add > 4 / length(cd_sample_model2_add))

```

```

large_cd_train = cd_sample_model2_add > 4 / length(cd_sample_model2_add)
plot(cd_sample_model2_add[large_cd_train])

coef(sample_model2)
sample_model2_add_fix = lm(trip_duration ~ dropoff_longitude,
                           data = train,
                           subset = cd_sample_model2_add <= 4 / length(cd_sample_model2_add))
coef(sample_model2_add_fix)

#set.seed(42)
#shapiro.test(resid(sample_model))

boxcox(sample_model2, plotit = TRUE)
attach(sub_train)
sample_model3 = lm(trip_duration ~
  pickup_datetime:dropoff_datetime+pickup_latitude:pickup_longitude+dropoff_latitude:dropoff_longitude+store_a
  family = negative.binomial(1))

par(mfrow = c(2,3))
plot(sample_model3,
      pch = 20,
      cex = 2,
      col = "grey")
abline(sample_model3, lwd = 3, col = "darkorange")
halfnorm(hatvalues(sample_model3))

summary(sample_model3)
#confint(sample_model3, level = 0.99)
tidy(sample_model3)
#augment(sample_model)
glance(sample_model3)
cooks.distance(sample_model3)[11] > 4 / length(cooks.distance(sample_model3))
bptest(sample_model3)
par(mfrow = c(2,3))
hist(resid(sample_model3))

qqnorm(resid(sample_model3), main = "Normal Q-Q Plot, sample_model", col = "darkgrey")
qqline(resid(sample_model3), col = "dodgerblue", lwd = 2)

plot(log(fitted(sample_model3)), log(resid(sample_model3)), col = "grey", pch = 20,
      xlab = "Fitted", ylab = "Residuals", main = "Data from Model 3")
abline(h = 0, col = "darkorange", lwd = 2)

plot(which(hatvalues(sample_model3) > 2 * mean(hatvalues(sample_model3)), TRUE))

plot(rstandard(sample_model3)[abs(rstandard(sample_model3)) > 2])
cd_sample_model3_add = cooks.distance(sample_model3)
sum(cd_sample_model3_add > 4 / length(cd_sample_model3_add))

large_cd_train = cd_sample_model3_add > 4 / length(cd_sample_model3_add)
plot(cd_sample_model3_add[large_cd_train])

coef(sample_model3)
sample_model3_add_fix = lm(trip_duration ~ dropoff_longitude,
                           data = train,
                           subset = cd_sample_model3_add <= 4 / length(cd_sample_model3_add))

```

```

coef(sample_model3_add_fix)

#set.seed(42)
#shapiro.test(resid(sample_model))

boxcox(sample_model3, plotit = TRUE)
m1AIC <- AIC(sample_model)
m1BIC <- BIC(sample_model)
m2AIC <- AIC(sample_model2)
m2BIC <- BIC(sample_model2)
m3AIC <- AIC(sample_model3)
m3BIC <- BIC(sample_model3)

AIC <- list(m1AIC, m2AIC, m3AIC)
BIC <- list(m1BIC, m2BIC, m3BIC)
kable(rbind(AIC, BIC), col.names = c("Model 1", "Model 2", "Model 3")) %>%
  kable_styling(full_width = T)
eval_p<-predict(sample_model3,sub_train, type = "response")
write.csv(eval_p,"predicted_eval_values.csv")
par(mfrow = c(1,3))
train_p<-predict(sample_model,sub_train, type = "response")
plot(train_p,sub_train$trip_duration,main = "Model 1")
train_p<-predict(sample_model2,sub_train, type = "response")
plot(train_p,sub_train$trip_duration,main = "Model 2")
train_p<-predict(sample_model3,sub_train, type = "response")
plot(train_p,sub_train$trip_duration,main = "Model 3")

```

## References:

1. Traffic congestion costs metro economy \$20 billion a year: study. Crain's New York Business. <https://www.craigslist.com/article/20180118/POLITICS/180119895/traffic-congestion-costs-metro-economy-20-billion-a-year-study>. Published January 18, 2018. Accessed November 13, 2019.
2. FDNY: Traffic — Not Bike Lanes — is to Blame for Increased Response Times. Streetsblog New York City. <https://nyc.streetsblog.org/2019/09/19/fdny-traffic-not-bike-lanes-is-to-blame-for-increase-d-response-times/>. Published September 19, 2019. Accessed November 13, 2019.
3. Skandul, E. How New York City taxis can get ahead of Uber and Lyft. CSNY. <https://www.cityandstateny.com/articles/opinion/commentary/how-new-york-city-taxis-can-get-ahead-of-uber-and-lyft.html>. Published October 17, 2019. Accessed November 13, 2019.
4. Sanders A. De Blasio struggles to find new Taxi and Limousine commissioner after first pick's disastrous performance at Council Hearing. nydailynews.com. <https://www.nydailynews.com/news/politics/ny-city-council-de-blasio-official-tlc-commissioner-jeff-lynch-20191003-cfx27yihafcebj4e62wytsprpu-story.html>. Published October 3, 2019. Accessed November 13, 2019.
5. New York City Taxi Trip Duration | Kaggle. Kaggle.com. <https://www.kaggle.com/c/nyc-taxi-trip-duration/data#>. Published 2017. Accessed November 13, 2019.
6. About TLC - TLC. Nyc.gov. <https://www1.nyc.gov/site/tlc/about/tlc-trip-record-data.page>. Published 2019. Accessed November 13, 2019.
7. Use holidays as a feature | Kaggle. Kaggle.com. <https://www.kaggle.com/c/nyc-taxi-trip-duration/discussion/37192#latest-572412>. Published 2017. Accessed November 13, 2019.