

Lecture 07  
2020 Spring Data-622  
LDA – basic Classifier 1934  
Raman Kannan

**Instructor Email Address:** [Raman.Kannan@sps.cuny.edu](mailto:Raman.Kannan@sps.cuny.edu)

Acknowledgements:  
Generous support from IBM Power Systems Academic Initiative  
IBM PSAI provides computing infrastructure for free

# Script for Algorithms

Develop the Intuition

Understand the assumptions

Develop the mathematics

Run the algorithms

Learn to interpret the result/output

Predict using the model

Learn to determine the performance

Distinguish training/testing error

Differentiate between overfitting/underfitting

Techniques to improve performance

# What are we doing?

## **We started with review**

- math/statistics/probability

We familiarized ourselves with

- Machine Learning

We took note of many tools

- required for ML

We practiced prediction

- quantitative predictors and response variables

We developed a process, performance metric

- RMSE in the case of regression

Then we encountered dichotomous response

- we transformed into logit
- logistic regression – probabilities
- Instead of RMSE we adopted AUC/Deviance

Are there other strategies to determine

- class (response) variable given predictor variables

This is an age old problem.

Qualitative response variables are all too common.

# Fisher's Linear Discriminant

A **dichotomy** /daɪˈkɒtəmi/ is a **partition** of a whole (or a set) into two parts (subsets). In other words, this couple of parts must be

- **jointly exhaustive**: everything must belong to one part or the other, and
- **mutually exclusive**: nothing can belong simultaneously to both parts.

Such a partition is also frequently called a bipartition.

<https://en.wikipedia.org/wiki/Dichotomy>

The original **dichotomous** discriminant analysis was developed by Sir **Ronald Fisher** in 1936.<sup>[8]</sup>

Almost 100 years ago

[https://en.wikipedia.org/wiki/Linear\\_discriminant\\_analysis](https://en.wikipedia.org/wiki/Linear_discriminant_analysis)

# LDA Formulation

If we have  $n$ -feature vectors, we can stack them into one matrix as follows;

$$Y = W^T X$$

$$\text{where } X_{m \times n} = \begin{bmatrix} x_1^1 & x_1^2 & \cdot & x_1^n \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ x_m^1 & x_m^2 & \cdot & x_m^n \end{bmatrix}, \quad Y_{C-1 \times n} = \begin{bmatrix} y_1^1 & y_1^2 & \cdot & y_1^n \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ y_{C-1}^1 & y_{C-1}^2 & \cdot & y_{C-1}^n \end{bmatrix}$$
$$\text{and } W_{m \times C-1} = [w_1 \mid w_2 \mid \dots \mid w_{C-1}]$$

If we compute  $W$ , given any new observation – using the above formula we can map it to the  $Y$  space...

# LDA Space

## 4 THE LDA SPACE

Linear Discriminant Analysis (LDA) [3], [4] searches for those vectors in the underlying space that best discriminate among classes (rather than those that best describe the data). More formally, given a number of independent features relative to which the data is described, LDA creates a linear combination of these which yields the largest mean differences between the desired classes. Mathematically speaking, for all the samples of all classes, we define two measures: 1) one is called *within-class* scatter matrix, as given by

$$S_w = \sum_{j=1}^c \sum_{i=1}^{N_j} (\mathbf{x}_i^j - \mu_j)(\mathbf{x}_i^j - \mu_j)^T,$$

where  $\mathbf{x}_i^j$  is the  $i$ th sample of class  $j$ ,  $\mu_j$  is the mean of class  $j$ ,  $c$  is the number of classes, and  $N_j$  the number of samples in class  $j$ ; and 2) the other is called *between-class* scatter matrix

$$S_b = \sum_{j=1}^c (\mu_j - \mu)(\mu_j - \mu)^T,$$

where  $\mu$  represents the mean of all classes.

Scatter  $S_W$  Within CLASS Scatter  
And  $S_B$  Between CLASS Scatter

$$S_W = \sum_{i=1}^c S_i$$

$$\text{where } S_i = \sum_{x \in \omega_i} (x - \mu_i)(x - \mu_i)^T$$

$$\text{and } \mu_i = \frac{1}{N_i} \sum_{x \in \omega_i} x$$

$$S_B = \sum_{i=1}^c N_i (\mu_i - \mu)(\mu_i - \mu)^T$$

$$\text{where } \mu = \frac{1}{N} \sum_{\forall x} x = \frac{1}{N} \sum_{\forall x} N_i \mu_i$$

$$\text{and } \mu_i = \frac{1}{N_i} \sum_{x \in \omega_i} x$$

# LDA Projection

- We can define the mean vectors for the projected samples  $\mathbf{y}$  as:

$$\tilde{\mu}_i = \frac{1}{N_i} \sum_{y \in \omega_i} y \quad \text{and} \quad \tilde{\mu} = \frac{1}{N} \sum_{\forall y} y$$

- While the scatter matrices for the projected samples  $\mathbf{y}$  will be:

$$\tilde{S}_W = \sum_{i=1}^C \tilde{S}_i = \sum_{i=1}^C \sum_{y \in \omega_i} (y - \tilde{\mu}_i)(y - \tilde{\mu}_i)^T$$

$$\tilde{S}_B = \sum_{i=1}^C N_i (\tilde{\mu}_i - \tilde{\mu})(\tilde{\mu}_i - \tilde{\mu})^T$$

LDA like PCA reduces dimension but preserves class semantics  
Also assumes a Gaussian (mean and variance) and all features are I.i.d  
Appears to work well if there is sufficient data otherwise performs < PCA  
Uses PCA

# LDA Implementation -1

```
pi_lda <- function(y){  
  pi_est ← table(y) / length(y)  
  return(as.matrix(pi_est))  
}
```

```
mu_lda <- function(X, y){  
  data_est <- as.data.frame(cbind(X,y))  
  data_est$X <- as.numeric(as.character(data_est$X))  
  mu <- aggregate(data = data_est, X ~ y, FUN = "mean")  
  colnames(mu) <- c("y", "X")  
  return(mu)  
}
```

```
var_lda <- function(X, y, mu){  
  n <- length(X)  
  K <- length(unique(y))  
  k <- unique(y)  
  var_est <- 0  
  
  for (i in 1:K){  
    var_est <- sum((X[y == k[i]] - mu$X[k[i]] == mu$y))^2) + var_est  
  }  
  
  var_est <- (1 / (n - K)) * var_est  
  return(var_est)  
}
```



# LDA Implementation -1.1

```
discriminant_lda <- function(X, pi, mu, var){  
  K <- length(unique(y))  
  k <- unique(y)  
  disc <- matrix(nrow = length(X), ncol = K)  
  colnames(disc) <- k  
  for (i in 1:K){  
    disc[,i] <- X * (mu$X[i] / var) - ((mu$X[i]^2) / (2 * var)) + log(pi[i])  
  }  
  
  disc <- as.data.frame(disc)  
  disc$predict <- apply(disc, 1, FUN = "which.max")  
  return(disc) }
```

```
X <- iris[,1]  
y <- as.character(iris[,5])  
  
pi_est <- pi_lda(y)  
mu_est <- mu_lda(X, y)  
var_est <- var_lda(X, y, mu_est)  
discriminant_est <- discriminant_lda(X, pi_est, mu_est, var_est)  
  
table(discriminant_est$predict, iris$Species)}
```

# LDA Performance

```
> discriminant_est <- discriminant_lda(X, pi_est, mu_est, var_est)
>
> table(discriminant_est$predict, iris$Species)
```

	setosa	versicolor	virginica
1	45	6	1
2	5	30	12
3	0	14	37

```
> |
```

```
> discriminant_est <- discriminant_lda(X, pi_est, mu_est, var_est)
> table(discriminant_est$predict, iris$Species)

      setosa versicolor virginica
1      45          6          1
2       5         30         12
3       0         14         37

> LDA_ConfMat<-table(discriminant_est$predict, iris$Species)
> print(paste("Accuracy is::",as.numeric(sum(diag(LDA_ConfMat))/sum(LDA_ConfMat)),sep=" "))
[1] "Accuracy is:: 0.7466666666666667"
> print(paste("Error is::",1-as.numeric(sum(diag(LDA_ConfMat))/sum(LDA_ConfMat)),sep=" "))
[1] "Error is:: 0.2533333333333333"
```

# Performance when more than 2 Classes

**Table 19.1** Confusion matrix for two classes.

True Class	Predicted class		Total
	Positive	Negative	
Positive	$tp$ : true positive	$fn$ : false negative	$p$
Negative	$fp$ : false positive	$tn$ : true negative	$n$
Total	$p'$	$n'$	$N$

As noted in Alpaydin's book most of these measures are not meaningful when there are more than 2 classes. We however computer Accuracy and Error.

**Table 19.2** Performance measures used in two-class problems.

Name	Formula
error	$(fp + fn)/N$
accuracy	$(tp + tn)/N = 1 - \text{error}$
tp-rate	$tp/p$
fp-rate	$fp/n$
precision	$tp/p'$
recall	$tp/p = \text{tp-rate}$
sensitivity	$tp/p = \text{tp-rate}$
specificity	$tn/n = 1 - \text{fp-rate}$

# LDA Let us clear the Smoke and Mirrors

Let us determine W and determine training error using Scatter matrices in R

```
head(iris)
table(iris$Species)
se<-iris[iris$Species=='setosa',]
ve<-iris[iris$Species=='versicolor',]
vi<-iris[iris$Species=='virginica',]
se<-iris[iris$Species=='setosa',]
semean<-apply(se[,1:4],2,mean)
vimean<-apply(vi[,1:4],2,mean)
vemean<-apply(ve[,1:4],2,mean)
allclassmean<-apply(iris[,1:4],2,mean)
(semean+vemean+vimean)/3 == allclassmean # confirm above numbers
S1<-cov(se[,1:4])
S2<-cov(vi[,1:4])
S3<-cov(ve[,1:4])
SW<-S1+S2+S3
N1<-nrow(se)
N2<-nrow(vi)
N3<-nrow(ve)
```

# LDA Let us clear the Smoke and Mirrors

Let us determine W and determine training error using Scatter matrices in R

```
SB1<-N1*(semean-allclassmean)*(semean-allclassmean)
SB2<-N2*(vimean-allclassmean)*(vimean-allclassmean)
SB3<-N3*(vemean-allclassmean)*(vemean-allclassmean)
SB<-SB1+SB2+SB3
invSW<-solve(SW)
#invSWSB<-invSW%*%SB
#DOES NOT WORK SWSBeigendecom<-eigen(invSWSB)
invSWSB<-invSW*SB
SWSBeigendecom<-eigen(invSWSB)
lambdas<-SWSBeigendecom$vectors
               [,which.max(SWSBeigendecom$values)]
as.matrix(iris[,1:4])%*% lambdas->iris.projections
rounded.iris.projections<-round(iris.projections)
```

*We computed the eigen values and used it to project data to Species.*

# LDA Let us clear the Smoke and Mirrors

```
iris.lda<-cbind(iris,classp=iris.projections,spectral=rounded.iris.projections)
iris.lda$predicted <-ifelse(iris.lda$spectral < (1-4.76),'virginica',
    ifelse(iris.lda$spectral>(1-3),'setosa','versicolor'))
iris.lda.confMat<-table(iris.lda$predicted,iris.lda$Species)
iris.lda.confMat
print(paste("Accuracy is::",
as.numeric(sum(diag(iris.lda.confMat))/sum(iris.lda.confMat)),sep= " "))
print(paste("Error is::",
1-as.numeric(sum(diag(iris.lda.confMat))/sum(iris.lda.confMat)),sep= " "))
```

```
> iris.lda.confMat<-table(iris.lda$predicted,iris.lda$Species)
> iris.lda.confMat

      setosa versicolor virginica
setosa      50         0         0
versicolor   0        25         0
virginica     0        25        50
> print(paste("Accuracy is::",as.numeric(sum(diag(iris.lda.confMat))/sum(iris.lda.confMat)),sep= " "))
[1] "Accuracy is:: 0.833333333333333"
> print(paste("Error is::",1-as.numeric(sum(diag(iris.lda.confMat))/sum(iris.lda.confMat)),sep= " "))
[1] "Error is:: 0.166666666666667"
```

We got better performance except that I personally set the threshold.

# QDA

LDA handles any type of data even if IV is not normally distributed.  
But LDA does assume all predictor variables have the same variance  
And the derivation relies on univariate Gaussian.

QDA allows multivariate scenario and uses multivariate gaussian.

## References:

R.A. Fisher, <sup>a</sup>The Statistical Utilization of Multiple Measurements,<sup>o</sup> Annals of Eugenics, vol. 8, pp. 376-386, 1938

<https://datascienceplus.com/how-to-perform-logistic-regression-lda-qda-in-r/>

<http://www2.ece.ohio-state.edu/~aleix/pami01.pdf>

[http://www.sci.utah.edu/~shireen/pdfs/tutorials/Elhabian\\_LDA09.pdf](http://www.sci.utah.edu/~shireen/pdfs/tutorials/Elhabian_LDA09.pdf)

[https://rstudio-pubs-](https://rstudio-pubs-static.s3.amazonaws.com/336635_7611ceab3e324623b9a7bea8de2b3818.html)

[static.s3.amazonaws.com/336635\\_7611ceab3e324623b9a7bea8de2b3818.html](https://rstudio-pubs-static.s3.amazonaws.com/336635_7611ceab3e324623b9a7bea8de2b3818.html)

# Well Known Package

```
require(MASS)
names(iris)
names(iris.lda)
mass_lda<-MASS::lda(Species~Sepal.Length+Sepal.Width+Petal.Length+Petal.Width,data=iris.lda)
mass_lda.predict<-predict(mass_lda,iris.lda[,1:4])
summary(mass_lda.predict)
str(mass_lda.predict)
mass_lda.predict
mass_lda.confmat <- table(mass_lda.predict$class,iris.lda$Species)
mass_lda.confmat
print(paste("Accuracy is::",as.numeric(sum(diag(mass_lda.confmat))/sum(mass_lda.confmat)),sep= " "))
print(paste("Error is::",1-as.numeric(sum(diag(mass_lda.confmat))/sum(mass_lda.confmat)),sep= " "))
history(25)
```

```
> mass_lda.confmat <- table(mass_lda.predict$class,iris.lda$Species)
> mass_lda.confmat
```

	setosa	versicolor	virginica
setosa	50	0	0
versicolor	0	48	1
virginica	0	2	49

```
>
> print(paste("Accuracy is::",as.numeric(sum(diag(mass_lda.confmat))/sum(mass_lda.confmat)),sep= " "))
[1] "Accuracy is:: 0.98"
> print(paste("Error is::",1-as.numeric(sum(diag(mass_lda.confmat))/sum(mass_lda.confmat)),sep= " "))
[1] "Error is:: 0.02"
```



# And the code for Visualizing is

```
> require(klaR)
Loading required package: klaR
Warning message:
package 'klaR' was built under R version 3.5.3
> partimat(Species~Sepal.Length+Sepal.Width+Petal.Length+Petal.Width,data=iris.lda,method="lda")
> |
```

require(klaR)

partimat(Species~Sepal.Length+Sepal.Width+Petal.Length+Petal.Width,data=iris.lda,method="lda")

# Visualize DecisionBoundaries

